Regime-mixture exceedance probabilities from a partly observed path, with collapse diagnostics

Moussa Ouallaf
Informatics Institute
University of Amsterdam

Abstract

We want the probability that a process ends above a threshold, given a path whose spread switches between levels. `regimix` fits a hidden Markov model (unobserved switching states) without labels to scaled recent increments and mixes per-state Gaussian tail probabilities by the filtered state weights; it fixes a variance-floor bug in hmmlearn 0.3.3. The finding is negative: the states beat a one-state model but not a plain volatility estimate. On 13,078 held-out EUR/USD 1-minute episodes, regimix ($K = 2$) minus arm in log-loss is -0.0079 nats for the one-state pooled probit (paired block-bootstrap SE 0.0014; negative favours regimix), +0.0024 (SE 0.0015, a tie) for a realised-volatility probit and +0.0043 (SE 0.0012, a loss) for a HAR-style probit; pristine hmmlearn EM is identical. On all 3 public datasets it ties or loses to the best volatility probit, and its outputs never improve gradient boosting (+0.0004 / +0.0011 / +0.0015 nats, SE 0.0004 / 0.0008 / 0.0004). SEs bootstrap time blocks or series of each dataset's last 30 %.

1 Problem

An episode i is a scalar process X observed up to a decision time t_i, a threshold k_i and a remaining horizon $\tau_i > 0$. The label is

$$y_i = \mathbf{1}\{X(t_i + \tau_i) > k_i\}, \tag{1}$$

with $\mathbf{1}\{\cdot\}$ the indicator of the event, and the goal is a calibrated $P(y_i = 1)$ from the path before t_i. The plain answer is a probit, a Gaussian tail probability whose width comes from the recent path, which works when volatility is stable but fails when activity jumps between a quiet and a busy level faster than a long realised-volatility window (the summed squared recent increments) can follow.

We expected a filtered regime, read off a few lookback windows, to react within one window and beat a slow volatility estimate. The filter is Hamilton's [Hamilton, 1989] and the regime-weighted normal tail is standard in regime-switching return models [Hardy, 2001]; we claim neither. The delta is how the parts are combined and checked: one label-free HMM fitted on the lookback windows of many short episodes, a calibration fitted on out-of-fold outputs, a variance-floor fix to the upstream library [hmmlearn developers, 2024], and diagnostics that say when the regimes are not real.

Preprint.

2 The change

Pick lookback edges $0 = e_0 < e_1 < \cdots < e_S$. Segment s is the increment over one window, scaled to unit time:

$$z_{i,s} = \mathrm{clip}\left(\frac{X(t_i - e_{s-1}) - X(t_i - e_s)}{\sqrt{e_s - e_{s-1}}},\, -c,\, c\right), \tag{2}$$

$$w_i = \alpha_{i,S}\, A, \qquad \alpha_{i,S}(j) = P(r_S = j \mid z_{i,1}, \ldots, z_{i,S}), \tag{3}$$

$$\tilde{p}_i = \sum_{j=1}^{K} w_i(j)\, \Phi\left(\frac{d_i + \mu_j \tau_i}{\max(\sigma_j, \sigma_{\min})\sqrt{\tau_i}}\right), \qquad d_i = X(t_i) - k_i, \tag{4}$$

$$p_i = \mathrm{sigmoid}\big(a + b\ \mathrm{logit}\, \tilde{p}_i\big). \tag{5}$$

Symbol	Meaning	Symbol	Meaning
$z_{i,s}$	segment s, oldest first	c	optional clip
r_s, K	regime at segment s; regime count	A	transition matrix
$\alpha_{i,S}$	filtered regime probabilities	w_i	horizon regime weights
μ_j, σ_j^2	drift, variance per unit time ($\mu_j = 0$)	$\sigma_{\min}$	spread floor (0.05)
Φ	normal CDF	a, b	Platt parameters [Platt, 1999, Lin et al., 2007]

We fit the hidden Markov model (HMM) with K Gaussian states by Baum–Welch [Baum et al., 1970] on the segment sequences of all episodes, without labels. Equation (3) is its forward pass, batched over episodes, moved one step ahead by A (next-step weighting), and Equation (4) averages one random-walk tail probability per regime. Labels enter only in Equation (5), through two parameters that we fit on outputs from 5 contiguous time-block folds before the HMM is refitted on all training rows. The estimator also returns w_i and $\mathrm{logit}\, \tilde{p}_i$, so that a downstream learner can use the regimes as inputs. We claim three properties, and Appendix A proves each of them.

Lemma 1 (exact mixture)**.** *If, given the horizon regime $r = j$, $X(t_i + \tau_i) - X(t_i) \sim \mathcal{N}(\mu_j \tau_i, \sigma_j^2 \tau_i)$ independently of $z_{i,1..S}$, $P(r = j \mid z_i) = w_i(j)$ and $\sigma_j \geq \sigma_{\min}$, then $\tilde{p}_i = P(y_i = 1 \mid z_i)$.*

Lemma 2 (no test label enters)**.** *If no training label reads a value of X inside the held-out period, then p_i for a held-out episode depends on labels only through training labels.*

The second part of the change is a variance floor: hmmlearn documents `min_covar` as a floor on the covariance but applies it only at initialisation, so on data with many exact zeros one state collapses onto a spike at zero. Its likelihood then grows without bound and its posterior takes only two values, a one-bit gate. Our `FlooredGaussianHMM` applies the floor after every M-step (the parameter-update half of EM):

$$v_j \leftarrow \max\big(v_j^{\mathrm{EM}},\, m\big)\,. \tag{6}$$

Here v_j^{EM} is the unconstrained variance update and m is `min_covar` (0.05 by default).

Lemma 3 (monotone EM)**.** *For diagonal or spherical covariances, update (6) is the exact maximiser of the M-step objective over $\{v \geq m\}$, so the log-likelihood never decreases.*

We turn the conditions under which regimes can help into warnings, computed from the fitted variances v_j and the self-transition probabilities A_{jj}:

$$\rho = \frac{\min_j v_j}{m}, \qquad D_j = \frac{1}{1 - A_{jj}}, \qquad R = \frac{\max_j \sigma_j}{\min_j \sigma_j}. \tag{7}$$

Here ρ is the collapse ratio (near 1: a state sits on the floor), D_j the mean dwell in segments and R the spread ratio. A dwell below 2 segments or a spread ratio below 2 triggers a warning; the report also gives the share of exactly flat segments, the count of distinct posterior values (a gate has two) and per-segment permutation importance. Appendix C gives the inputs, the four calls that apply regimix to any path, and the commands that reproduce every number here.

3 Setup

Data. We use 3 public datasets, each chosen because its activity level switches. *EUR/USD 1-minute* bars from HistData, 2016–2023: at a random minute each hour, $P(\text{price in 30 min} > \text{price} + u)$ with

$u \in \{-3, 0, +3\}$ pips (1 pip $= 0.0001$ in price). *US large-cap daily* closes for 59 tickers, 2005–2024, one decision day in 4 per ticker: $P(\text{close in 10 days} > \text{close} \times (1 + u))$, $u \in \{-2, 0, +2\}\,\%$. ***M4 hourly***, the hourly series of the M4 forecasting competition [Makridakis et al., 2020], one decision in 9 hours: $P(\log y_{t+h} > \log y_t + u)$ for $h \in \{12, 24\}$ hours and $u = c\,\hat{\sigma}_1\sqrt{h}$ ($\hat{\sigma}_1$ the one-step training spread), $c \in \{-0.5, 0, 0.5\}$. Appendix B lists the lookback edges. The held-out set is the last 30 % of time (per series for M4) and is never shuffled; we drop every training row whose label reads into it, an embargo (a gap against look-ahead) of one horizon. That leaves 34,202 / 13,078, 51,360 / 22,125 and 51,952 / 23,288 training / held-out episodes, each series scaled by the spread of its one-step training increments.

Arms. All arms see the same training rows, segments, distances and horizons, and we give each a Platt map fitted on training rows. No arm gets a hyperparameter search: each runs fixed settings written down before the run. *regimix* runs its defaults ($K = 2$, `n_iter` 60, `tol` 0.01, floor and next-step weighting $w_i = \alpha_{i,S}A$ on, drift off), and the *one-regime pooled probit*, our trivial baseline, is the same model with $K = 1$: one constant spread. The *realised-vol probit* uses the realised variance per unit time over the lookback; the *HAR-style probit* (heterogeneous autoregressive volatility) regresses log realised variance over the horizon on its value at each lookback window [Corsi, 2009, Andersen et al., 2003]. The *gate ablation* replaces w_i by a one-hot weight on the high-spread state when $|z_{i,S}|$ exceeds the training quantile that matches that state's stationary share (its long-run fraction of segments under A). *Logistic regression* and *histogram gradient boosting* [Pedregosa et al., 2011] get $(Z, |Z|, d, \tau, d/\sqrt{\tau})$. *Pristine hmmlearn EM* is the upstream 0.3.3 M-step [hmmlearn developers, 2024] without the floor (our batched E-step matches upstream to 10^{-14}), and the stacking test [Wolpert, 1992] appends out-of-fold w_i and $\operatorname{logit} \tilde{p}_i$ to the boosting inputs.

Metric and error. Log-loss in nats, averaged over 3 seeds; the paired difference per episode is

$$\Delta_i = \ell(p_i^{\text{regimix}}, y_i) - \ell(p_i^{\text{arm}}, y_i), \qquad \ell(p, y) = -y \log p - (1 - y) \log(1 - p). \tag{8}$$

Negative $\bar{\Delta}$ means regimix is better. Episodes in one time block share a path, so a per-episode SE is too small (FX against $K = 1$: 0.0011 against 0.0014). We therefore sum Δ_i within blocks and run a moving-block bootstrap [Künsch, 1989] (resampling runs of consecutive blocks) with 2,000 draws over 749 FX days and 1,500 equity dates, with the block length set by the first lag at which the autocorrelation of block means falls inside $2/\sqrt{n}$. M4 series are independent, so there we resample 414 whole series (a Newey–West SE [Newey and West, 1987] is kept as a check). We call a difference a tie when its 95 % interval covers zero. Calibration is the expected calibration error over 15 bins (ECE, the mean gap between predicted and observed frequency) [Guo et al., 2017]. The 15 arms of the main run were written down before any held-out number existed. After the main run lost to the volatility probits we tried 6 post-hoc variants on the same held-out period, about 18 comparisons, and kept a variant only if it won on all 3 datasets.

4 Results

No dataset gives regimix a win over the best volatility probit (Table 1). On EUR/USD the regimes pass both diagnostics: they persist 19.6 and 8.7 segments on average and their spreads differ 3.37-fold. There regimix beats the pooled probit by 0.0079 nats and the gate ablation by 0.0075. The realised-vol probit is a tie (+0.0024, interval -0.0005 to +0.0053) and the HAR probit wins by 0.0043 (+0.0019 to +0.0067). Our reading, not tested, is that two discrete states are a coarse copy of the continuous volatility level that the HAR arm estimates directly.

On equities (Figure 1) the regimes pass both warnings (dwell 14.9 / 3.9, spread ratio 2.72), yet every probit lands within 0.0008 nats of regimix. On M4 the diagnostics warn on both counts (dwell 1.0 / 1.3 segments, spread ratio 1.09), and M4 is exactly where regimix loses to $K = 1$, by 0.0029 nats. We count its edge over the HAR probit there (-0.0007, interval -0.0013 to 0.0000) as a tie. On M4 logistic regression and boosting reach 0.3476 and 0.2126, far below every probit, which we attribute, without a test, to daily seasonality and mean reversion that a driftless random walk cannot express; M4 therefore ranks regimix only among probits.

Appended to boosting, the regime outputs change log-loss by +0.0004 and +0.0011 (ties) and by +0.0015 on M4, where the interval +0.0007 to +0.0023 excludes zero on the wrong side.

Table 1: Regimix wins against the pooled probit only where regimes exist, and never against the best volatility probit. Held-out log-loss, nats; in brackets regimix minus arm ± bootstrap SE (negative: regimix better). Last row: boosting with regime outputs minus boosting (positive: appending hurt).

arm	EUR/USD 1-min (n = 13,078)	US large-cap daily (n = 22,125)	M4 hourly (n = 23,288)
regimix ($K = 2$)	**0.6201**	**0.6505**	**0.5878**
pooled probit ($K = 1$)	0.6280 (-0.0079 ± 0.0014)	0.6508 (-0.0003 ± 0.0005)	0.5849 (+0.0029 ± 0.0002)
realised-vol probit	0.6177 (+0.0024 ± 0.0015)	0.6505 (0.0000 ± 0.0008)	0.5869 (+0.0010 ± 0.0006)
HAR-style probit	0.6158 (+0.0043 ± 0.0012)	0.6496 (+0.0008 ± 0.0005)	0.5885 (-0.0007 ± 0.0003)
gate ablation	0.6276 (-0.0075 ± 0.0012)	0.6541 (-0.0036 ± 0.0005)	0.5910 (-0.0032 ± 0.0003)
logistic regression	0.6288 (-0.0087 ± 0.0016)	0.6523 (-0.0018 ± 0.0008)	0.3476 (+0.2402 ± 0.0049)
gradient boosting	0.6195 (+0.0006 ± 0.0010)	0.6530 (-0.0025 ± 0.0008)	0.2126 (+0.3752 ± 0.0053)
pristine hmmlearn 0.3.3 EM	0.6201 (0.0000 ± 0.0000)	0.6505 (0.0000 ± 0.0000)	0.5878 (0.0000 ± 0.0000)
boosting + regime − boosting	+0.0004 ± 0.0004	+0.0011 ± 0.0008	+0.0015 ± 0.0004

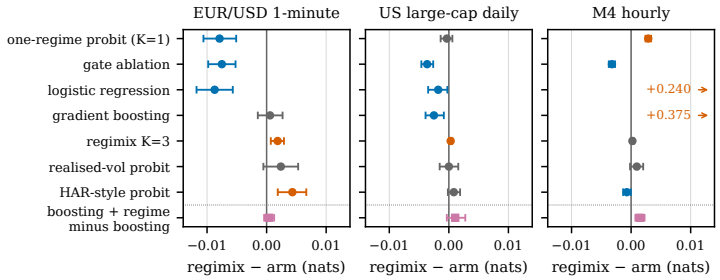


Figure 1: The volatility probits match or beat regimix on every dataset, and the regime outputs never help boosting. Paired held-out log-loss difference, regimix minus arm, 95 % block or (M4) series bootstrap intervals; blue a regimix win, red a loss, grey a tie, purple stacking.

Calibration follows the same order (Figure 2; ECE gaps untested): on EUR/USD the seed-averaged ECE is 0.0181 for regimix, 0.0237 for the pooled probit and 0.0122 for HAR, while on M4 regimix has the highest ECE of the three, 0.0623.

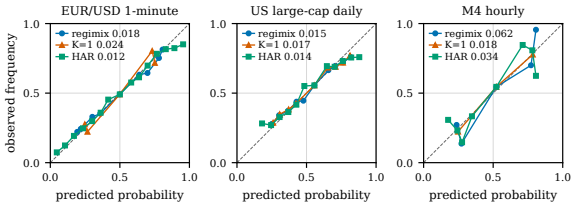


Figure 2: Regimix has a lower ECE than the pooled probit where regimes exist and a higher one on M4, where they do not. Reliability diagrams of the seed-averaged held-out probabilities, 15 bins, sparse bins hidden; the legend gives each arm's ECE.

Of the single-default ablations (Table 2, Appendix B) only $K = 3$ helps, still short of the volatility probits; we keep $K = 2$ because $K = 3$ was picked on the held-out period, and no post-hoc variant won on all 3 datasets. Out-of-fold calibration, ablated after the main run, is no clear gain over in-sample Platt (regimix minus it: 0.0000, -0.0004, +0.0008), while dropping calibration costs most on M4 (-0.0149 ± 0.0009).

5 Limits

The idea lost a comparison and stays open: a realised-volatility probit, with no hidden states and the same two calibration parameters, reads the current volatility level as well as the states do. To beat it, the states must carry what the level misses, such as a switch inside the horizon [Hardy, 2001], which regimix does not model. The evidence is thin: one held-out period per dataset, 3 seeds, post-hoc variants on that period, a floor with only synthetic evidence, survivorship bias in the equities (only today's large caps), a drift unit mismatch (Appendix B), and no ablation of fitting one HMM on all episodes.

References

Torben G. Andersen, Tim Bollerslev, Francis X. Diebold, and Paul Labys. Modeling and forecasting realized volatility. *Econometrica*, 71(2):579–625, 2003. doi: 10.1111/1468-0262.00418.

Leonard E. Baum, Ted Petrie, George Soules, and Norman Weiss. A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains. *The Annals of Mathematical Statistics*, 41(1):164–171, 1970. doi: 10.1214/aoms/1177697196.

Fulvio Corsi. A simple approximate long-memory model of realized volatility. *Journal of Financial Econometrics*, 7(2):174–196, 2009. doi: 10.1093/jjfinec/nbp001.

Arthur P. Dempster, Nan M. Laird, and Donald B. Rubin. Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society: Series B (Methodological)*, 39(1): 1–22, 1977. doi: 10.1111/j.2517-6161.1977.tb01600.x.

Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. On calibration of modern neural networks. In *Proceedings of the 34th International Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Research*, pages 1321–1330, 2017. URL https://proceedings.mlr.press/v70/guo17a.html.

James D. Hamilton. A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2):357–384, 1989. doi: 10.2307/1912559.

Mary R. Hardy. A regime-switching model of long-term stock returns. *North American Actuarial Journal*, 5(2):41–53, 2001. doi: 10.1080/10920277.2001.10595984.

hmmlearn developers. hmmlearn: Hidden Markov models in Python, with scikit-learn like API. https://github.com/hmmlearn/hmmlearn, 2024. Version 0.3.3.

Hans R. Künsch. The jackknife and the bootstrap for general stationary observations. *The Annals of Statistics*, 17(3):1217–1241, 1989. doi: 10.1214/aos/1176347265.

Hsuan-Tien Lin, Chih-Jen Lin, and Ruby C. Weng. A note on Platt's probabilistic outputs for support vector machines. *Machine Learning*, 68(3):267–276, 2007. doi: 10.1007/s10994-007-5018-6.

Spyros Makridakis, Evangelos Spiliotis, and Vassilios Assimakopoulos. The M4 competition: 100,000 time series and 61 forecasting methods. *International Journal of Forecasting*, 36(1):54–74, 2020. doi: 10.1016/j.ijforecast.2019.04.014.

Whitney K. Newey and Kenneth D. West. A simple, positive semi-definite, heteroskedasticity and autocorrelation consistent covariance matrix. *Econometrica*, 55(3):703–708, 1987. doi: 10.2307/1913610.

Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011. URL https://jmlr.org/papers/v12/pedregosa11a.html.

John C. Platt. Probabilistic outputs for support vector machines and comparisons to regularized likelihood methods. In Alexander J. Smola, Peter Bartlett, Bernhard Schölkopf, and Dale Schuurmans, editors, *Advances in Large Margin Classifiers*, pages 61–74. MIT Press, 1999.

David H. Wolpert. Stacked generalization. *Neural Networks*, 5(2):241–259, 1992. doi: 10.1016/S0893-6080(05)80023-1.

A Proofs

Lemma 1. **Given:** 1) law of total probability over a partition $\{r = j\}_{j=1}^{K}$; 2) given $r = j$, the increment $U = X(t_i + \tau_i) - X(t_i) \sim \mathcal{N}(\mu_j \tau_i, \sigma_j^2 \tau_i)$, independent of z_i; 3) $P(r = j \mid z_i) = w_i(j)$; 4) $\Phi(-x) = 1 - \Phi(x)$.

Choose: the event $E = \{y_i = 1\} = \{X(t_i) + U > k_i\} = \{U > -d_i\}$ and the partition of rule 1.

Need to prove: $P(E \mid z_i) = \sum_j w_i(j)\, \Phi\big((d_i + \mu_j \tau_i)/(\sigma_j \sqrt{\tau_i})\big)$.

$$
\begin{aligned}
P(E \mid z_i) &= \textstyle\sum_j P(E \mid r = j, z_i)\, P(r = j \mid z_i) && \text{rule 1} \\
&= \textstyle\sum_j P(E \mid r = j)\, w_i(j) && \text{independence in rule 2, rule 3} \\
P(E \mid r = j) &= P\Big(\tfrac{U - \mu_j \tau_i}{\sigma_j \sqrt{\tau_i}} > \tfrac{-d_i - \mu_j \tau_i}{\sigma_j \sqrt{\tau_i}}\Big) && \text{standardise } U\text{, rule 2} \\
&= 1 - \Phi\Big(\tfrac{-d_i - \mu_j \tau_i}{\sigma_j \sqrt{\tau_i}}\Big) && \text{definition of } \Phi \\
&= \Phi\Big(\tfrac{d_i + \mu_j \tau_i}{\sigma_j \sqrt{\tau_i}}\Big) && \text{rule 4}
\end{aligned}
$$

Precondition: Equation (4) uses $\max(\sigma_j, \sigma_{\min})$, which equals σ_j because $\sigma_j \geq \sigma_{\min}$. **Therefore,** substitution gives $\tilde{p}_i = P(y_i = 1 \mid z_i)$. □

Lemma 2. **Given:** 1) $\theta = \mathrm{fit}(Z_{\text{train}})$, Baum–Welch on segments only; 2) $(a, b) = \mathrm{Platt}\big(\{\mathrm{logit}\, \tilde{p}_j^{\text{oof}}, y_j\}_{j \in \text{train}}\big)$, where each out-of-fold output comes from an HMM fitted on segments only; 3) for a held-out episode, $p_i = g(z_i, d_i, \tau_i; \theta, a, b)$ by Equations (3)–(5).

Choose: a held-out episode i and write p_i through the fitted quantities of rules 1 and 2.

Need to prove: p_i depends on labels only through $\{y_j\}_{j \in \text{train}}$.

$$
\begin{aligned}
p_i &= g(z_i, d_i, \tau_i;\, \theta,\, a,\, b) && \text{rule 3} \\
&= g\big(z_i, d_i, \tau_i;\, \mathrm{fit}(Z_{\text{train}}),\, \mathrm{Platt}(\{\mathrm{logit}\, \tilde{p}_j^{\text{oof}}, y_j\}_{j \in \text{train}})\big) && \text{rules 1 and 2}
\end{aligned}
$$

The only label symbols on the right are y_j with $j \in$ train.

Precondition: a training label must not itself carry held-out information. A training label reads $X(t_j + \tau_j)$; rows with $t_j + \tau_j$ inside the held-out period are dropped (the embargo of Section 3), so every remaining y_j is a function of X before the held-out period.

Therefore, no held-out label enters p_i. The out-of-fold step is not needed for this; it only makes the calibrator see raw outputs from an HMM that did not fit those rows, as at test time. □

Lemma 3. **Given:** 1) generalised EM [Dempster et al., 1977]: if each update satisfies $Q(\theta^{\text{new}} \mid \theta^{\text{old}}) \geq Q(\theta^{\text{old}} \mid \theta^{\text{old}})$, the log-likelihood does not decrease; 2) for diagonal or spherical covariances Q separates into one term per state and coordinate, and the variance term has the form $q(v) = -\frac{1}{2}(\nu \log v + s/v)$ with weight $\nu > 0$ and scatter $s > 0$ (a conjugate variance prior only changes ν and s); 3) the mean update does not depend on v.

Choose: one variance term q and the feasible set $v \geq m$.

Need to prove: $\arg\max_{v \geq m} q(v) = \max(s/\nu,\, m)$.

$$
\begin{aligned}
q'(v) &= -\tfrac{1}{2}\left(\tfrac{\nu}{v} - \tfrac{s}{v^2}\right) && \text{derivative of rule 2} \\
&= \tfrac{\nu}{2v^2}\left(\tfrac{s}{\nu} - v\right) && \text{factor out } \tfrac{\nu}{2v^2} > 0
\end{aligned}
$$

So $q' > 0$ for $v < s/\nu$ and $q' < 0$ for $v > s/\nu$: q is unimodal with peak $v^* = s/\nu = v^{\text{EM}}$. Same method on each side of m. **Case** $v^* \geq m$**:**

$$
\begin{aligned}
\arg\max_{v \geq m} q(v) &= v^* && \text{peak is feasible, unimodal } q \\
&= \max(v^*, m) && v^* \geq m
\end{aligned}
$$

Case $v^* < m$**:**

$$
\begin{aligned}
\arg\max_{v \geq m} q(v) &= m && q' < 0 \text{ on } [m, \infty) \text{ since } m > v^* \\
&= \max(v^*, m) && v^* < m
\end{aligned}
$$

Substitution gives update (6) in both cases.

Precondition: rule 1 needs θ^{old} feasible, so that it is a candidate of the constrained maximisation. The initial variances carry $+m$, so they are feasible. Every later variance is the output of (6), so it is $\geq m$. By induction every iterate is feasible.

Therefore, by rule 3 the means keep their usual update, each variance gets its exact constrained maximiser, Q does not decrease, and by rule 1 neither does the log-likelihood. □

B Ablations and post-hoc variants

The benchmark arm equals the public `RegimeExceedanceClassifier` with a maximum logit difference of 0. The lookback edges are 0, 1, 3, 7, 15, 31, 63 minutes (EUR/USD), 0, 1, 3, 7, 15, 31 days (equities) and 0, 1, 2, 4, 8, 16, 32, 48 hours (M4). The tie on M4 for $K = 3$ has interval -0.0001 to +0.0005. Turning drift on is a large loss on equities and M4; behind `drift=True` the segment mean is drift per $\sqrt{\text{time}}$ while Equation (4) reads it per unit time. Neither floor binds on the public data: the smallest stock-EM variance is 4.8, 8.3 and 30.5 times the floor, so both ablations are bit-identical. The floor's evidence is synthetic. On 4,000 sequences of length 7 drawn as rounded $\mathcal{N}(0, 0.5^2)$ values (about 68 % exact zeros), stock `GaussianHMM` with $m = 10^{-3}$ ends with variances 0.163 and 2.3×10^{-6}, which is 0.002 times the floor.

Table 2: No single default carries the method, the floors never bind on public data, and calibration matters most where regimes are not real (M4). Regimix minus the ablated arm, log-loss in nats ± paired SE; a positive value means the ablated arm is better.

ablation	EUR/USD 1-min	US large-cap daily	M4 hourly
$K = 3$	+0.0019 ± 0.0006	+0.0003 ± 0.0001	+0.0002 ± 0.0001
pristine EM (no floor)	0.0000 ± 0.0000	0.0000 ± 0.0000	0.0000 ± 0.0000
next step off	+0.0001 ± 0.0002	-0.0005 ± 0.0001	-0.0003 ± 0.0000
$\sigma_{\min}$ off	0.0000 ± 0.0000	0.0000 ± 0.0000	0.0000 ± 0.0000
drift on	0.0000 ± 0.0001	-0.0093 ± 0.0028	-0.1028 ± 0.0024
isotonic calibration	0.0000 ± 0.0004	-0.0003 ± 0.0003	+0.0076 ± 0.0006
per-regime Platt	+0.0002 ± 0.0003	-0.0010 ± 0.0002	+0.0020 ± 0.0003
in-sample Platt	0.0000 ± 0.0000	-0.0004 ± 0.0001	+0.0008 ± 0.0000
no calibration	+0.0008 ± 0.0004	-0.0052 ± 0.0032	-0.0149 ± 0.0009

All 6 variants were tried after the main run, on the same held-out period. Values are variant minus the named arm, ± paired SE (negative: the variant is better). The level-relative variants divide segments and distance by the realised-vol or HAR level, so the regimes model only the change around it.

Table 3: Dividing by the volatility level turns regimix into a near-copy of the level arm, and more regimes do not help against it. No variant wins on all three datasets.

variant	against	EUR/USD 1-min	US large-cap daily	M4 hourly
$K = 4$	realised vol	-0.0003 ± 0.0012	+0.0001 ± 0.0008	+0.0031 ± 0.0007
$K = 4$	HAR	+0.0016 ± 0.0010	+0.0009 ± 0.0005	+0.0015 ± 0.0006
rv-relative, $K = 2$	realised vol	-0.0002 ± 0.0001	0.0000 ± 0.0000	+0.0022 ± 0.0003
rv-relative, $K = 3$	realised vol	-0.0005 ± 0.0002	0.0000 ± 0.0000	+0.0060 ± 0.0005
HAR-relative, $K = 2$	HAR	0.0000 ± 0.0001	0.0000 ± 0.0000	+0.0022 ± 0.0002
HAR-relative, $K = 3$	HAR	-0.0001 ± 0.0000	0.0000 ± 0.0000	+0.0011 ± 0.0002
$n = 300$	realised vol	+0.0035 ± 0.0015	+0.0051 ± 0.0011	-0.0041 ± 0.0007
$n = 1{,}000$	realised vol	+0.0021 ± 0.0015	-0.0016 ± 0.0007	-0.0026 ± 0.0006
$n = 3{,}000$	realised vol	+0.0032 ± 0.0017	-0.0004 ± 0.0008	-0.0028 ± 0.0006
$n = 300$	boosting	-0.0281 ± 0.0023	-0.1556 ± 0.0087	+0.1881 ± 0.0046

The only result we have outside these benchmarks follows. On a proprietary binary threshold-exceedance task, the regime-mixture output did not beat a single-regime realised-volatility probit on its own: block-CV log-loss +0.010 nats worse (paired block-bootstrap SE 0.012, 95 % CI [-0.013, +0.032], i.e. a tie within noise). Against stock hmmlearn on held-out log-loss, a zero-inflated emission gave +0.0001 (SE 0.0037), a per-position variance +0.0014 (SE 0.0012), and a Student-t emission -0.0040 (SE 0.0015, 95 % CI [-0.0068, -0.0009]; negative is better).

The small-n rows train on the most recent 300, 1,000 or 3,000 episodes. With few labels regimix beats boosting by a wide margin on EUR/USD and equities, but not the realised-vol probit. M4 training rows are ordered by series, so its small-n rows test transfer to other series, not a learning curve.

C Applying regimix to new data, and reproduction

Inputs. One observed path (or many), the decision times, a threshold per decision and a horizon. Divide the path by the spread of its one-step increments on the training part, so that segments and distances are in unit scale. Pick lookback edges that roughly double, for example $0, 1, 3, 7, 15, 31$ steps, with the oldest covering a few horizons. Keep episodes in time order: the out-of-fold calibration needs a splitter of contiguous time blocks, and the held-out set must come after the training set with an embargo of one horizon.

Calls. `Z = segments_from_path(x, times, t, edges)` builds Equation (2); `clf = RegimeExceedanceClassifier(n_states=2).fit(Z, y, distance=d, horizon=h)` fits the HMM and the calibrator, with $d = X(t) - k$; `clf.predict_proba(Z, distance=d, horizon=h)` and `clf.regime_posterior(Z)` return p_i and w_i; `diagnostics.report(clf, Z, y=y, distance=d, horizon=h)` prints the warnings of Equation (7). If the report warns on dwell or spread, the regimes are not real and $K = 1$ or a realised-vol probit is the better choice, as on M4. Always run a realised-vol probit beside it: on our data it was the arm to beat. `examples/quickstart.py` runs all of this on synthetic data with a known answer.

Reproduction. Package, benchmark and paper live in the `regimix` repository; the results in `bench/results/` are from commit `a6e9ca4`. The whole delta to upstream is `patches/hmmlearn/0001-gaussianhmm-apply-min-covar-after-every-m-step.patch` (and a spherical-shape fix, `0002`) against hmmlearn 0.3.3, with repro tests in `patches/hmmlearn/tests`. `bash bench/run_all.sh` downloads the data and writes `bench/results/{fx,eq,m4}[_variants].json` (about 1 hour on 2 threads); `python -m paper.results.ablate_calib {fx,eq,m4}` writes the calibration ablation to `paper/results/calib_*.json`, `python results/make_numbers.py` regenerates every number in this paper and `python figures/make.py` both figures (Figure 2 reads the held-out predictions the benchmark stores outside git).